

# Counting visitors with wuepix

*Jeroen Staab*

*2018-10-29*

The *wuepix* package counts visitor numbers using computer vision. Therefore three methods (Change Detection, HOG-Descriptor, YOLO-darknet) were wrapped into this package. Additional management tools, as a Ground-Truth-Data sampler, are also included here. This vignette demonstrates a typical workflow.

## Packages

```
# Installation
# devtools::install_github("georoen/wuepix")
library(wuepix)
library(tidyverse)
```

## Site configuration

### Paths & Filenames

Define the directory paths and filename patterns implied by the data archive.

```
# Where to find Images?
## Raw data
img.folder_raw <- "IMG_raw/"
# Preprocessed (cropped, scaled, enhanced,...)
img.folder <- "IMG/"
# Remove corrupted images by filesize (in byte)
threshold <- 10000
# How to grep date?
gsub.Date <- function(Filename){gsub("picam-", "", gsub(".jpg", "", Filename))}
# Date code
date.code <- "%Y%m%d-%H%M"

# Aggregation Scale
T_scale <- "20 mins"

# Hubland
# threshold <- 1000
# gsub.Date <- function(Filename){gsub("Camera1_M_", "", gsub(".jpg", "", Filename))}
# date.code <- "%Y-%m-%d-%H%M%S"
```

### Extent of interest

To speed up processing an extend of interest (EOI) should be selected. Using the Linux comandline tool *ImageMagick*, this can also include rotations as well as other image operations. However identifying the correct command involves visual interpretation of the results. To do so I proceeded as follows.

## 1. Using Gimp / Photoshop

Initially use *GIMP* to identify the pre-process routine (bounding-box, optional rotation).

Tipp: Overlay several images to cover different scenarios.

## 2. Test comandline

After identifying the pre-process routine try to put the parameters into *ImageMagick* and test command on a single image using *convert*.

```
# Test
# SizeX x SizeY + PostionX + PositionY
convert.string <- "-crop 1600x800+0+1030"
cmd <- paste("convert extra/Ref_raw.jpg", convert.string, "extra/Ref.jpg")
system(cmd)
message("Please check cropped Ref.jpg, then proceed")
```

This results in the following extend of interest. Only this part of the image will be further analysed, so please only proceed if satisfied with the result.

## 3. Preprocess image archive

Next all images will be pre-processed according to the routine developed above using *mogrify*. Please pay attention to the slightly different syntax of the *mogrify -crop ... -path IMG/ IMG\_raw/\*.jpg*. This will preprocess all images from *IMG\_raw/* and save them in *IMG/*.

```
# Preprocess
dir.create(img.folder)
cmd <- paste("mogrify", convert.string, "-path", img.folder,
            paste0("\",img.folder_raw, \"*.jpg\", \"\""))
system(cmd)
message("Finished preprocessing")
```

## List images

First all images need to be listed. The following chunk does so, plus enhances the data frame according to *Site configuration*: (1) due to external effects (eg. transmission) images can be corrupted. Here files with a file size smaller than the *threshold* will be exluded. (2) The Timestamp gets interpreted, therefore first the filenames are cropped with help of *gsub.Date*. Because filenames can be very different and the corresponding regular expression can very complex, it seemed easiest to do with a function. This also makes developing it more simple due better testing option. After cropping the timestamp it will be converted to a *POSIXlt* time object using *date.code*. (3) Last but not least the relative filepaths are reconstructed. Note, that this should also work with *list.files(..., fullnames=TRUE)* but I remeber then struggling with grepping the datecode.

```
Files <- data.frame(Filename=list.files(img.folder, pattern = "*.jpg"),
                    stringsAsFactors = FALSE)

# Remove corrupted images
Files$Size <- file.size(paste0(img.folder, Files$Filename)) > threshold
Files <- Files[which(Files$Size),]
Files <- select(Files, -Size)
```

```

# Add Timestamp
Files$Timestamp <- strptime(gsub.Date(Files$Filename), date.code)
Files$Timestamp <- as.POSIXct(Files$Timestamp)
Files <- Files[order(Files$Timestamp),] # Order by Timestamp

# Full Filename
Files$Filename <- paste0(img.folder, Files$Filename)

```

To get an overview about the data beeing processed, here some metadata summarys are prompted.

```

## 21 files to analyze
## Dates from 24.06.2017 12:30 to 24.06.2017 12:50
## Time difference of 20 mins

```

## Ground-Truth-Data

To latter assess the classifiers accuracies, Ground-Truth-Data is mandatory. Use `GTD_list()` to manually count pedestrians in `Files$Filename`. Here all images (100%) got evaled, for sampling uncomment orange lines.

```

start <- Sys.time() # Get start time
#GTD <- GTD_list(sample(x = Files$Filename, size = 10))
#the.sample <- sample(c(1:nrow(Files)), size = 100)
#Files <- Files[the.sample,]
Files$GTD <- GTD_list(Files$Filename)
Files$GTD <- as.numeric(Files$GTD)
(Sys.time() - start) # Print runtime

save(Files, file = "Results/GTD.RData")
write.csv(Files, file = "Results/GTD.csv")

```

Sum of visitors in GTD.

To aggregate the time-series by `T_scale` use `fun_Aggregation()`. Use tidyverse grammar to select wished aggeration method (mean or sum).

```

# Aggregation
Files_res <- fun_Aggregation(Files$Timestamp, Files$GTD, T_scale) %>%
  select(-MEAN) %>%
  rename(GTD = SUM)

```

## Processing

Finally we can start processing the (preprocessed) image archive. For a detailed description of the methods, the reader is refered to the authors master thesis.

### Method 1: Change detection

This approach is inspired by methods used remote sensing and biotech. Using algebra, two pictures are applied against each other, revealing changes. Use `CD_list()` for processing a list of images, including parallel processing. See `?CD_single()` for the available parameters.

```

# Processing
start <- Sys.time() # Get start time
CD <- CD_list(Files$Filename, Min = 0.9, method = "ratio",
              predictions = "CD_Predictions")
(Sys.time() - start) # Print runtime

## Time difference of 7.604359 secs

Files$Hum <- CD
# Aggregation
Files_res <- fun_Aggregation(Files$Timestamp, Files$Hum, T_scale) %>%
  select(-SUM) %>%
  rename(CD = MEAN) %>%
  left_join(Files_res)

## Joining, by = "Timestamp"

Now convert the number of changed pixels into visitor numbers by calibrating them.

# Calibration
lm_cal <- lm(GTD ~ 0+CD, data = Files_res)
summary(lm_cal)

##
## Call:
## lm(formula = GTD ~ 0 + CD, data = Files_res)
##
## Residuals:
##      1      2
## -3.155  3.442
##
## Coefficients:
##      Estimate Std. Error t value Pr(>|t|)
## CD  0.07309    0.03516   2.079   0.285
##
## Residual standard error: 4.669 on 1 degrees of freedom
## Multiple R-squared:  0.8121, Adjusted R-squared:  0.6242
## F-statistic: 4.322 on 1 and 1 DF,  p-value: 0.2854

Files_res$CD_pred <- round(predict(lm_cal, select(Files_res, -GTD)))

```

## Method 2: Histogramms of Oriented Gradients

The second method uses *Histogramms of Oriented Gradients* to detect pedestrians [DALAL\_2005]. The HOG-Descriptor, focuses on a feature class of the same name. Here the *OpenCV*-Python implementation was wrapped into `hog_list()`. For an installation guide please see `?hog_install()`. First however upscale the images according to the trainingset.

```

# Resize
dir.create("IMG_resize/")

## Warning in dir.create("IMG_resize/"): 'IMG_resize' existiert bereits

cmd <- paste("mogrify -resize 270x240 -path IMG_resize/",
             paste0("IMG/", "*.jpg"))
system(cmd)

```

```

message("Finished preprocessing")

## Finished preprocessing
Files_resized <- gsub("IMG/", "IMG_resize/", Files$Filename)

# Processing
start <- Sys.time() # Get start time
HOG <- hog_list(Files_resized, resize = 1, padding = 24, winStride = 2,
               Mscale = 1.05, predictions = "HOG_Predictions/")
(Sys.time() - start) # Print runtime

## Time difference of 1.489887 secs

To access the object-based accuracies, run GTD_truePositives(). It will cor.test() the input for you and
returns the miss-rate, FPPW and so forth.

Files$HOG <- HOG
GTD_truePositives(Files$GTD, Files$HOG)

##
## Pearson's product-moment correlation
##
## data: GTD and PRD
## t = 9.7321, df = 19, p-value = 8.136e-09
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.7936218 0.9643845
## sample estimates:
## cor
## 0.9126413

## FN TP FP MR FPPW cor
## 1 6 8 0 0.4285714 0 0.9126413

# Aggregation
Files_res <- fun_Aggregation(Files$Timestamp, Files$HOG, T_scale) %>%
  select(-MEAN) %>%
  rename(HOG = SUM) %>%
  left_join(Files_res)

## Joining, by = "Timestamp"

```

## Method 3: Convolutional-Neural-Network

The third and last object detector discussed in this study bases on Convolutional-Neural-Networks. Here *YOLO* [REDMON\_2016] was utilized, as it comes pre-trained and has a open license. To install it use `yolo_install()`. Accordingly it was wrapped into `yolo_list()` and `yolo_single()`, whereby unfortunately only the latter is capable of saving the predictions. Using `sapply()` is possible, however note, that processing does take much longer, as the weights need to get loaded repeatedly. Attention, after every *wuepix* installation it is necessary to run `yolo_update()` as well, as this write a small file into the package installation linking to the *YOLO* installation. This links aids running the wrapper functions conveniently.

```

# Processing
start <- Sys.time() # Get start time
YOLO <- yolo_list(Files_resized)

```

```
# YOLO <- sapply(Files_resized, yolo_single, predictions = "YOLO_Predictions")
(Sys.time() - start) # Print runtime
```

```
## Time difference of 8.530758 mins
```

You can also use `GTD_truePositives()` here.

```
Files$YOLO <- YOLO
GTD_truePositives(Files$GTD, Files$YOLO)
```

```
##
## Pearson's product-moment correlation
##
## data: GTD and PRD
## t = 58.938, df = 19, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.9931527 0.9989179
## sample estimates:
## cor
## 0.9972763

## FN TP FP MR FPPW cor
## 1 1 13 0 0.07142857 0 0.9972763
```

```
# Aggregation
Files_res <- fun_Aggregation(Files$Timestamp, Files$YOLO, T_scale) %>%
  select(-MEAN) %>%
  rename(YOLO = SUM) %>%
  left_join(Files_res)
```

```
## Joining, by = "Timestamp"
```

As YOLO detects a lot of objects, `yolo_list()` logs a complete list of all detections to a separate file `yolo_detections.txt`. To read it in a tidy way use `yolo_Read()`.

```
# Read, group and count detections
yolo.results <- yolo_Read("yolo_detections.txt")
unique(yolo.results$Class)
```

```
## [1] "backpack" "bicycle" "person" "umbrella" "sports ball"
```

## Save Results

Last but not least, save the *R* environment for further studies.

```
save.image(file = "Results/Enviroment.RData")
```